

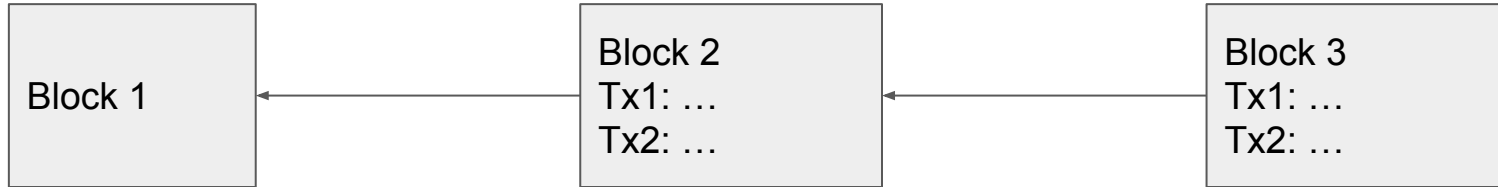
Erasure Coded Sharding

Bribery-resistant sharding for scalable blockchains

Jessica Taylor (jessi.liu.taylor@gmail.com)

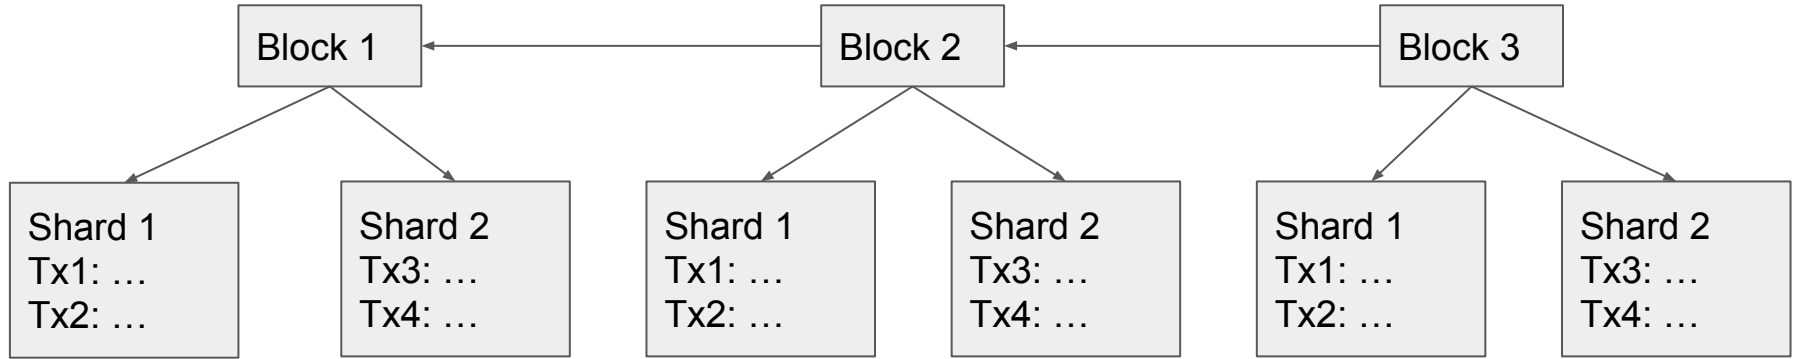
gigascaling.net

Regular blockchains (Bitcoin, Ethereum)



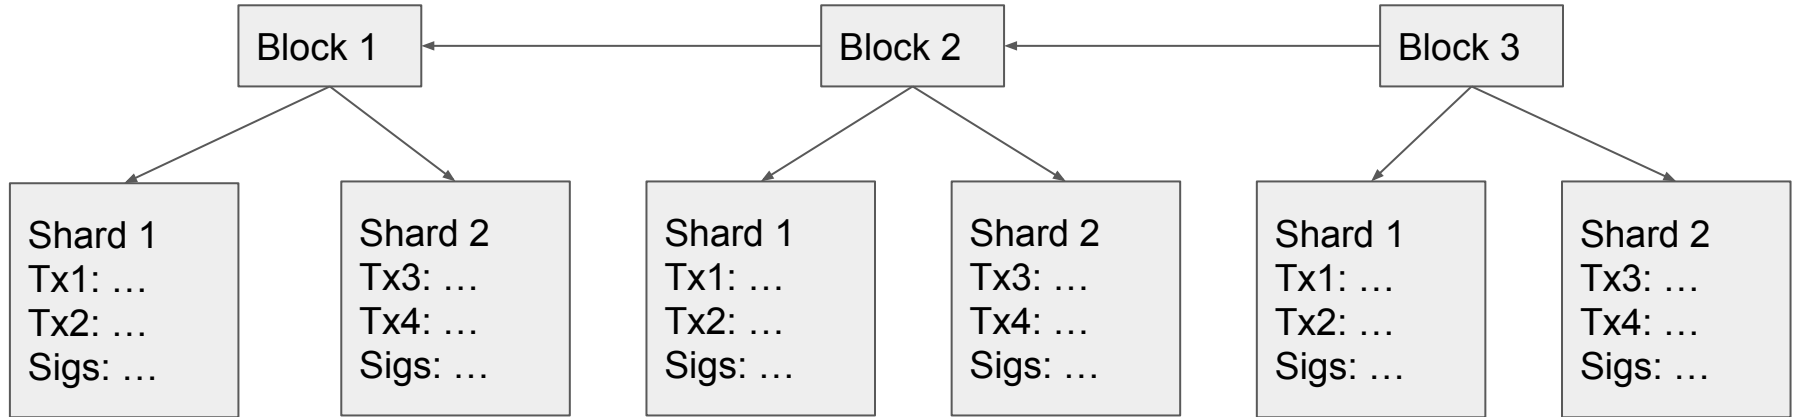
- “ $A \rightarrow B$ ” means “A includes hash code of B”
- Proof-of-work or proof-of-stake (details omitted)
- All full nodes download and verify every transaction
- Problem: this limits how many transactions per second there can be

Sharded blockchain






- Regular full nodes only download top chain
- Shard-level verifiers download data for a given shard
- Each shard is responsible for a subset of accounts
- Split transactions into “send” and “receive” (like actor model)
- Problem: how do regular full nodes know shard-level data is valid?

Random validator pools

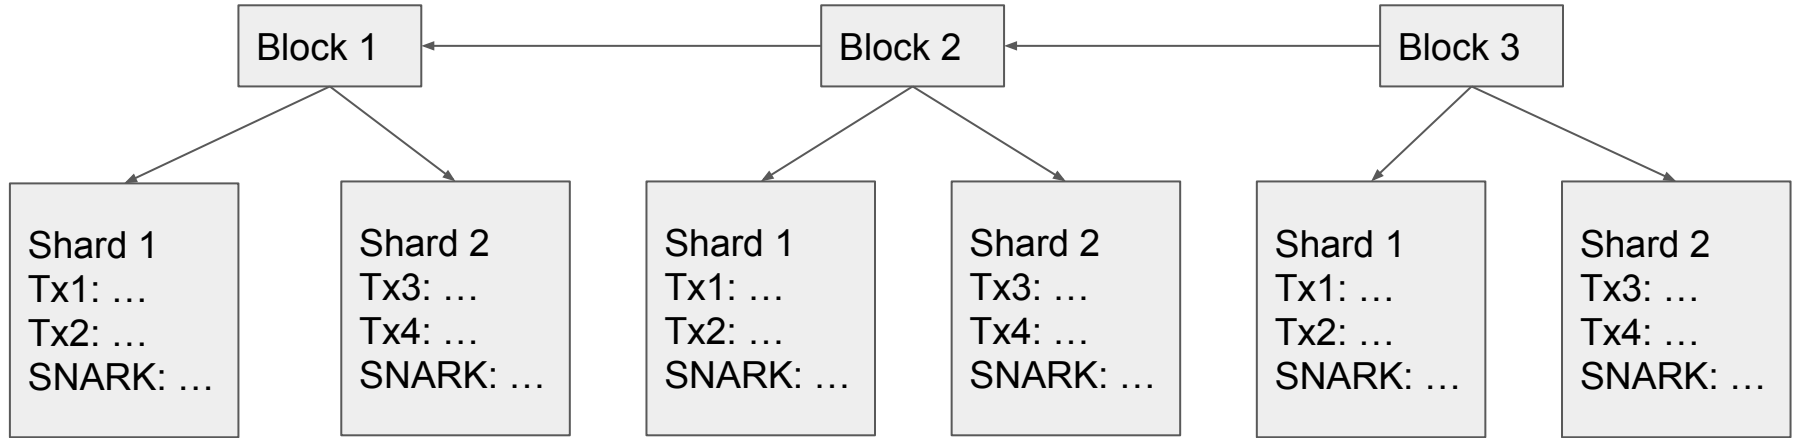


- Call together a validator pool by selecting ~1000 random stakers
- If $>\frac{2}{3}$ of stakers are honest, with high probability most of the pool is honest
- Regular full nodes only download signatures, not transactions
- Can be iterated multiple levels (Inductive Consensus Tree Protocol, ictp.io)
- Problem: bribery

3 types of validators

-  Altruistic: Always honest
-  Malicious: Trying to sabotage the system, even at cost to self
-  Greedy: Honest unless bribed sufficiently
- If $> \frac{1}{3}$ of stakers are malicious or greedy, quorums are likely to be bribable.
- Only one quorum needs to be bribed to mint fake currency; bribery is cheap.

zk-SNARK verification




- Each valid shard contains a zero-knowledge proof of its validity
- Full nodes download and check SNARKs, not transactions
- Can be iterated multiple levels (and compressed) with recursive SNARKs
- Problem: data availability attacks (reveal SNARK, withhold data)



Data availability attacks

- Shard-level validators may provide a SNARK showing the shard is available, but withhold the shard data (incl. transactions) from others
- This prevents individual accounts from proving account states or unspent transactions, which prevents the account's funds from being accessible
- Can we use random validator pools to mitigate this? They can be bribed, knocking out lots of account data
- Solution concept: Can we force a large *percentage* of pools to be knocked out to make *any* shard's data unavailable?
- (inspired by [PolyShard](#), algorithm details differ)





Reed-Solomon code

- 
- k blue squares = original data chunks (each the same # bytes), 5 in this case
- Redundancy factor α , 3 in this case
- $k(\alpha - 1)$ green squares = data augmentation
- Can recover k blue squares given *any* k blue or green squares!
- Works by polynomial interpolation: interpret blue squares as polynomial coefficients in a finite field, get green squares by evaluating polynomial at more points, recover blue squares by fitting a polynomial
- Encoding is $\sim O(\alpha k^2)$, decoding is $\sim O(\alpha^2 k^3)$
- Parallelizes well, tractable to do in a SNARK





Reed-Solomon coding shard data

- Shard 1 data: 
- Shard 2 data: 
- Etc, for x shards
- The data consists of $k\alpha$ chunks for each shard ($xk\alpha$ total chunks)
- Split into $k\alpha$ columns; we can recover all data from k full columns
- Split stakers into $k\alpha$ equal-sized pools, each responsible for storing 1 column
- If enough honest stakers store all their data, we can recover everything!
- Storage per staker = x chunks ($1/k$ of original data), it scales decently
- Use SNARKs to prove that enough signatures exist that a significant fraction of stakers would have to lie about their storage for data to be unrecoverable
- In case some chunks are dropped, use a different SNARK for each shard

Reducing storage by Reed-Solomon coding each column

- Shard 1 data: 
- Shard 2 data: 
- Redundant 1: 
- Redundant 2: 
- Secondary redundancy factor β (2 in this case)
- We can recover a column with x of βx chunks
- We can recover everything with $(\alpha + \beta - 1)xk$ of $\alpha\beta xk$ chunks
- Split stakers into $\alpha\beta xk$ equal-sized pools, each responsible for storing 1 chunk
- If enough honest stakers store their chunk, we can recover everything!
- Storage per staker = 1 chunk ($1/(xk)$ of original data), it scales very well

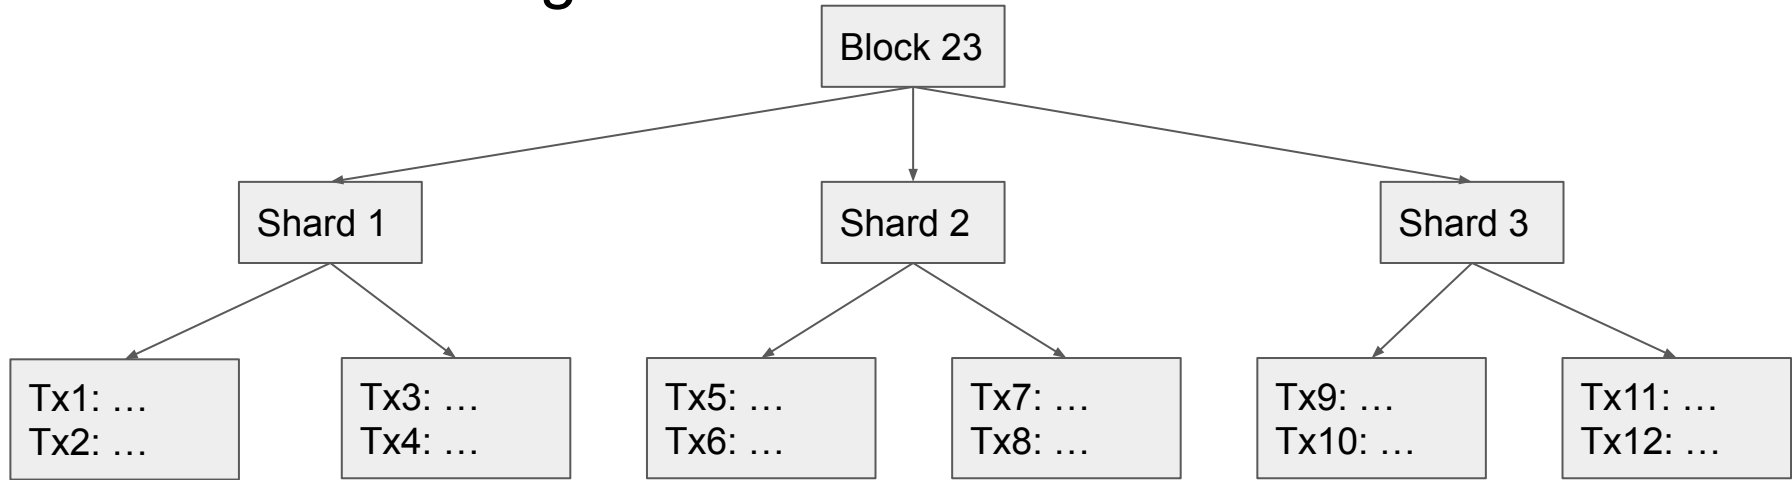
Security analysis

- Set threshold γ , proportion of chunks that must be *asserted* to be stored
- Proportion that must *actually* be stored is $(\alpha + \beta - 1)/(\alpha\beta)$
- Secure if less than $\gamma - (\alpha + \beta - 1)/(\alpha\beta)$ proportion of stakers act maliciously
- Malicious includes  and bribed 
- Cost of bribery \approx stake amount (easy to prove chunk unavailability & punish)
- Functional if more than γ proportion of stakers are honest most rounds
- Acting honestly includes  and  (bribery won't happen most rounds)
- E.g. $\gamma = \frac{2}{3}$, $\alpha = \beta = 6 \rightarrow$ secure if less than 36.1% of stakers act maliciously

Privacy and smart contracts

- Privacy is easy since we're already using SNARKs
- Transaction data for a private transaction consists of account ID and hash code of new private state (64 bytes total)
- SNARK for a shard proves SNARKs exist for each account state transaction
- 2 types of smart contracts: per-account and independent
- Per-account smart contracts modify account data, including private data
- A single account can partake in multiple per-account smart contracts
- E.g. tokens which are similarly private to the base currency
- Independent contracts are like their own account, data is public

Multi-level sharding



- 3 instead of 2 levels, reduced branching factor
- Reduces work per node
- Use 2d Reed-Solomon code for each shard, 3d overall (decreases data efficiency, increases compute efficiency of encoding/decoding)